

muhumuza Stephen

March 24, 2024

```
[1]: # importing necessary libraries
import numpy as np
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, PolynomialFeatures
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import matplotlib.pyplot as plt
```

```
[2]: # importing my dataset
data = pd.read_csv('C:
↳\\Users\\lynda\\Desktop\\Assignment\\student_scores_dataset.csv')
data.head()
```

```
[2]:
```

	Study Hours	Exam Scores
0	3.7	87.9
1	9.5	143.6
2	7.3	123.7
3	6.0	99.9
4	1.6	64.5

```
[3]: # identifying my variables
x=np.array(data['Study Hours']).reshape(-1,1)
x
```

```
[3]: array([[3.7],
           [9.5],
           [7.3],
           [6. ],
           [1.6],
           [1.6],
           [0.6],
           [8.7],
           [6. ],
           [7.1],
           [0.2],
           [9.7],
           [8.3],
```

[2.1],
[1.8],
[1.8],
[3.],
[5.2],
[4.3],
[2.9],
[6.1],
[1.4],
[2.9],
[3.7],
[4.6],
[7.9],
[2.],
[5.1],
[5.9],
[0.5],
[6.1],
[1.7],
[0.7],
[9.5],
[9.7],
[8.1],
[3.],
[1.],
[6.8],
[4.4],
[1.2],
[5.],
[0.3],
[9.1],
[2.6],
[6.6],
[3.1],
[5.2],
[5.5],
[1.8],
[9.7],
[7.8],
[9.4],
[8.9],
[6.],
[9.2],
[0.9],
[2.],
[0.5],
[3.3],

```

[3.9],
[2.7],
[8.3],
[3.6],
[2.8],
[5.4],
[1.4],
[8. ],
[0.7],
[9.9],
[7.7],
[2. ],
[0.1],
[8.2],
[7.1],
[7.3],
[7.7],
[0.7],
[3.6],
[1.2],
[8.6],
[6.2],
[3.3],
[0.6],
[3.1],
[3.3],
[7.3],
[6.4],
[8.9],
[4.7],
[1.2],
[7.1],
[7.6],
[5.6],
[7.7],
[4.9],
[5.2],
[4.3],
[0.3],
[1.1]])

```

```

[4]: # my dependent variable
y=data['Exam Scores']
y.head()

```

```

[4]: 0      87.9
     1     143.6

```

```
2    123.7
3     99.9
4     64.5
Name: Exam Scores, dtype: float64
```


```
[5]: # checking for missing values
data.isnull().sum
```

```
[5]: <bound method NDFrame._add_numeric_operations.<locals>.sum of      Study Hours
Exam Scores
0         False      False
1         False      False
2         False      False
3         False      False
4         False      False
..          ...        ...
95        False      False
96        False      False
97        False      False
98        False      False
99        False      False
```

```
[100 rows x 2 columns]>
```

```
[6]: #visualising the data
plt.scatter(x,y, label='Data_points',alpha=0.6, color='blue', s=75)
plt.title('Scatter plot of Study hours against exams score')
plt.xlabel('Study Hours')
plt.ylabel('Exams score')
```

```
[6]: Text(0, 0.5, 'Exams score')
```



```
[7]: # data preproccing
x_train,x_test,y_train,y_test = train_test_split(x,y, test_size=0.2,
↪random_state=42)
```

```
[8]: #x-train
x_train
```

```
[8]: array([[9.2],
          [8.9],
          [2. ],
          [0.3],
          [9.9],
          [1.8],
          [1.2],
          [5.2],
          [7.1],
          [0.1],
          [9.7],
          [5.2],
          [3.3],
```

[5.9],
[5.6],
[1.6],
[1.4],
[5.4],
[8.1],
[3.],
[1.8],
[9.7],
[8.7],
[4.9],
[5.1],
[2.9],
[6.2],
[7.9],
[8.3],
[2.1],
[4.6],
[6.],
[5.2],
[6.8],
[6.],
[3.6],
[0.6],
[2.8],
[3.],
[4.7],
[0.9],
[1.1],
[6.],
[9.1],
[9.7],
[8.],
[3.1],
[0.7],
[2.7],
[4.3],
[1.2],
[5.],
[0.5],
[5.5],
[0.3],
[2.],
[7.3],
[0.7],
[7.7],
[3.3],

```
[3.6],  
[3.1],  
[1. ],  
[0.5],  
[9.5],  
[9.4],  
[1.4],  
[7.3],  
[3.7],  
[6.4],  
[7.1],  
[7.1],  
[7.3],  
[3.3],  
[6.1],  
[3.9],  
[2. ],  
[1.8],  
[7.6],  
[7.8]])
```

```
[9]: #x_test  
x_test
```

```
[9]: array([[0.6],  
[8.9],  
[7.7],  
[6.6],  
[2.6],  
[4.4],  
[2.9],  
[8.6],  
[0.2],  
[3.7],  
[4.3],  
[6.1],  
[8.2],  
[9.5],  
[1.2],  
[1.6],  
[7.7],  
[0.7],  
[8.3],  
[1.7]])
```

```
[10]: # standardazing the independent variable  
scaler= StandardScaler()
```

```
x_train_scaled=scaler.fit_transform(x_train)
x_test_scaled=scaler.transform(x_test)
```

```
[11]: #creating model or building modele
model= LinearRegression()
model.fit(x_train_scaled,y_train)
```

```
[11]: LinearRegression()
```

```
[12]: #Evaluating the model
y_pred=model.predict(x_test_scaled)
mae= mean_absolute_error(y_test,y_pred)
mse= mean_squared_error(y_test,y_pred)
r2=r2_score(y_test,y_pred)

print('Mean Absolute Error:',mae)
print('Mean Squared Error:',mse)
print('R-squared:',r2)
```

Mean Absolute Error: 2.9365732667749755
Mean Squared Error: 16.202109700645348
R-squared: 0.9826924926918468

```
[13]: # interpreting my coefficients
model.intercept_
print('model.intercept',model.intercept_)
print('The intercept is 96.5875. It means that if the study hours are zero, the_
↪predicted exam score would be 96.5875.')
```

model.intercept_ 96.5875
The intercept is 96.5875. It means that if the study hours are zero, the predicted exam score would be 96.5875.

```
[14]: model.coef_
print('model.coef',model.coef_)
print(' the coefficient for study hours is 28.52556103. It means that for every_
↪one-unit increase in study hours, the exam scores are expected to increase_
↪by approximately 28.53, assuming all other factors remain constant.')
```

model.coef [28.52556103]
the coefficient for study hours is 28.52556103. It means that for every one-unit increase in study hours, the exam scores are expected to increase by approximately 28.53, assuming all other factors remain constant.

```
[15]: #model improvement
#Applying polynomial Features
poly= PolynomialFeatures(degree=2)
x_train_poly =poly.fit_transform(x_train)
x_test_poly =poly.transform(x_test)
```



```
[16]: #scaling
scaler = StandardScaler()
x_train_scaled1=scaler.fit_transform(x_train_poly)
x_test_scaled1 = scaler.transform(x_test_poly)
```

```
[17]: #training the linearRegression model with polynomial features
model_poly = LinearRegression()
model_poly.fit(x_train_scaled1,y_train)
```

```
[17]: LinearRegression()
```

```
[18]: #making prediction
y_pred_poly= model_poly.predict(x_test_scaled1)

y_pred_poly
```

```
[18]: array([ 57.67108889, 138.71725719, 126.05217671, 114.72389663,
          75.79896173,  92.87476622,  78.59488185, 135.52095869,
          54.15228189,  86.1485391 ,  91.90720366, 109.66365139,
          131.2903679 , 145.16991094,  63.01602912,  66.62380909,
          126.05217671,  58.55635145, 132.34467911,  67.5313149 ])
```

```
[19]: #Evaluation of my mode
mae_poly= mean_absolute_error(y_test,y_pred_poly)
mse_poly = mean_squared_error(y_test,y_pred_poly)
r2_poly = r2_score(y_test,y_pred_poly)
print('Mean Absolute error:',mae_poly,)
print('Mean Squared Error (Polynomial):',mse_poly)
print('R-squared (polynomial):',r2_poly)
```

Mean Absolute error: 2.828595846925279

Mean Squared Error (Polynomial): 15.643638436299161

R-squared (polynomial): 0.9832890659571593

```
[ ]: #comparison in evaluation
```

Based on these comparisons, it appears that the model **with** polynomial features **outperforms** the initial model **in** terms of accuracy, performance, **and** fit to the data.

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```